

# Phase 3 Machine Learning Project- SYRIATEL CUSTOMER CHURN

- Student : Spencer Lugalia
- Student Pace : DSPT08
- Instructor: Samuel Karu, Daniel Ekale

## 1. Business understanding

The dataset pertains to SyriaTel telecommunications company seeking to understand and reduce customer churn. Churn, defined as customers leaving the service, significantly impacts revenue and profitability. Key business objectives include identifying factors contributing to churn, such as call usage patterns, subscription plans, and customer interactions, to predict at-risk customers. This analysis aims to uncover actionable insights to improve customer retention strategies, enhance service offerings, and focus on critical features like international plan subscriptions, daytime call usage, and customer service interactions to reduce churn rates and improve customer satisfaction.

### 1.1 Business problem

SyriaTel communication company is in need of a classifier to predict whether a customer will ("soon") stop doing business with SyriaTel, a telecommunications company. This is a binary classification problem. Detecting customers who are likely to leave and putting retention plans in place to keep them are the primary business challenges. The company wants to reduce revenue loss by lowering turnover of customers.

### 1.2 The scope

The scope of the project include:

- Establishing if there are any predictable patterns whether a customer will ("soon") stop doing business with SyriaTel.
- Develop a predictive model to determine whether a customer will churn (binary classification: Yes/No) based on customer usage patterns, interaction with the company, and plan features.
- Provide actionable insights to SyriaTel to reduce customer churn by identifying high-risk customers and enabling targeted retention strategies.

### 1.3 Data Source

The Churn in Telecom's dataset was sourced from [Churn Data](#)

## 2 Data Understanding

Data understanding lets us explore and analyze our churn data to gain insights into its structure, content, and relationships. It involves looking at the types of data, identifying patterns, checking for missing values, and understanding the distribution of variables.

```
In [421]: # importing relevant libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
import seaborn as sns  
from sklearn.model_selection import train_test_split, cross_validate  
from sklearn.linear_model import LogisticRegression  
from sklearn.preprocessing import OneHotEncoder, StandardScaler  
from sklearn.compose import ColumnTransformer  
from sklearn.pipeline import Pipeline  
from sklearn.preprocessing import MinMaxScaler  
from sklearn.preprocessing import OneHotEncoder  
from sklearn.tree import DecisionTreeClassifier  
from sklearn import tree  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score  
from sklearn.metrics import accuracy_score, precision_score, recall_score  
from imblearn.over_sampling import SMOTE  
from sklearn.model_selection import GridSearchCV  
from sklearn.model_selection import RandomizedSearchCV
```

### **Load dataset**

In [422]:

# Read data from csv file,Checking the first 5 rows.  
data = pd.read\_csv('./Data/bigml\_59c28831336c6604c800002a.csv')  
data.head()

Out[422]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34

5 rows × 21 columns



In [423]:

#Checking the different columns info includig dtype and null counts  
data.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3333 entries, 0 to 3332  
Data columns (total 21 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   state                                3333 non-null   object  
1   account length                       3333 non-null   int64  
2   area code                           3333 non-null   int64  
3   phone number                         3333 non-null   object  
4   international plan                   3333 non-null   object  
5   voice mail plan                      3333 non-null   object  
6   number vmail messages                3333 non-null   int64  
7   total day minutes                    3333 non-null   float64  
8   total day calls                      3333 non-null   int64  
9   total day charge                     3333 non-null   float64  
10  total eve minutes                    3333 non-null   float64  
11  total eve calls                      3333 non-null   int64  
12  total eve charge                     3333 non-null   float64  
13  total night minutes                  3333 non-null   float64  
14  total night calls                    3333 non-null   int64  
15  total night charge                   3333 non-null   float64  
16  total intl minutes                   3333 non-null   float64  
17  total intl calls                     3333 non-null   int64  
18  total intl charge                    3333 non-null   float64  
19  customer service calls               3333 non-null   int64  
20  churn                               3333 non-null   bool  
dtypes: bool(1), float64(8), int64(8), object(4)  
memory usage: 524.2+ KB
```

- The data has 0 null values
- The data has various data types ranging from Object,integers, Floats and Booleans
- The structure of the data is 21 columns and 3333 entries

```
In [424]: ▶ # checking on all the column names
data.columns
```

```
Out[424]: Index(['state', 'account length', 'area code', 'phone number',
                'international plan', 'voice mail plan', 'number vmail messages',
                'total day minutes', 'total day calls', 'total day charge',
                'total eve minutes', 'total eve calls', 'total eve charge',
                'total night minutes', 'total night calls', 'total night charge',
                'total intl minutes', 'total intl calls', 'total intl charge',
                'customer service calls', 'churn'],
                dtype='object')
```

#### Columns overview

1. **state**: The state in which the customer resides.
2. **account length**: The duration (in months) the customer has had an account with the service provider.
3. **area code**: The area code associated with the customer's phone number.
4. **phone number**: The customer's phone number.
5. **international plan**: A binary indicator (yes/no) of whether the customer has an international calling plan.
6. **voice mail plan**: A binary indicator (yes/no) of whether the customer has a voicemail plan.
7. **number vmail messages**: The total number of voicemail messages received by the customer.
8. **total day minutes**: The total number of minutes the customer spent on daytime calls.
9. **total day calls**: The total number of daytime calls made by the customer.
10. **total day charge**: The total charge for daytime calls made by the customer.
11. **total eve minutes**: The total number of minutes the customer spent on evening calls.
12. **total eve calls**: The total number of evening calls made by the customer.
13. **total eve charge**: The total charge for evening calls made by the customer.
14. **total night minutes**: The total number of minutes the customer spent on nighttime calls.
15. **total night calls**: The total number of nighttime calls made by the customer.
16. **total night charge**: The total charge for nighttime calls made by the customer.
17. **total intl minutes**: The total number of minutes the customer spent on international calls.
18. **total intl calls**: The total number of international calls made by the customer.
19. **total intl charge**: The total charge for international calls made by the customer.
20. **customer service calls**: The total number of calls the customer made to customer service.
21. **churn**: A binary indicator (1/0) representing whether the customer has churned (left the service) or not.

In [425]: `# Descriptive statistics for the numerical columns in the dataset`  
`data.describe().T`

Out[425]:

	count	mean	std	min	25%	50%	75%	max
<b>account length</b>	3333.0	101.064806	39.822106	1.00	74.00	101.00	127.00	243.00
<b>area code</b>	3333.0	437.182418	42.371290	408.00	408.00	415.00	510.00	510.00
<b>number vmail messages</b>	3333.0	8.099010	13.688365	0.00	0.00	0.00	20.00	51.00
<b>total day minutes</b>	3333.0	179.775098	54.467389	0.00	143.70	179.40	216.40	350.80
<b>total day calls</b>	3333.0	100.435644	20.069084	0.00	87.00	101.00	114.00	165.00
<b>total day charge</b>	3333.0	30.562307	9.259435	0.00	24.43	30.50	36.79	59.64
<b>total eve minutes</b>	3333.0	200.980348	50.713844	0.00	166.60	201.40	235.30	363.70
<b>total eve calls</b>	3333.0	100.114311	19.922625	0.00	87.00	100.00	114.00	170.00
<b>total eve charge</b>	3333.0	17.083540	4.310668	0.00	14.16	17.12	20.00	30.91
<b>total night minutes</b>	3333.0	200.872037	50.573847	23.20	167.00	201.20	235.30	395.00
<b>total night calls</b>	3333.0	100.107711	19.568609	33.00	87.00	100.00	113.00	175.00
<b>total night charge</b>	3333.0	9.039325	2.275873	1.04	7.52	9.05	10.59	17.77
<b>total intl minutes</b>	3333.0	10.237294	2.791840	0.00	8.50	10.30	12.10	20.00
<b>total intl calls</b>	3333.0	4.479448	2.461214	0.00	3.00	4.00	6.00	20.00
<b>total intl charge</b>	3333.0	2.764581	0.753773	0.00	2.30	2.78	3.27	5.40
<b>customer service calls</b>	3333.0	1.562856	1.315491	0.00	1.00	1.00	2.00	9.00

In [426]: `#checking the shape of the data`  
`data.shape`  
`print(f"This data set consists of {data.shape[0]} rows")`  
`print(f"This data set consists of {data.shape[1]} columns")`

This data set consists of 3333 rows  
 This data set consists of 21 columns

In [427]: `# Confirming there are no Null values`  
`data.isnull().values.any()`

Out[427]: False

In [428]: `# Checking the total number of duplicated rows`  
`data.duplicated().sum()`  
`print(f"This data set consists of duplicated {data.duplicated().sum()} rows")`


This data set consists of duplicated 0 rows

## 2.1 Data Cleaning and Feature Engineering

At this point we have already established that there is 0 null and 0 duplicated rows therefore we will not have to fill null values or drop duplicated rows.

Therefore for this section i will do the following tasks:

- For the model drop non critical columns
- One-Hot Encode the 3 categorical columns 'international plan', 'voice mail plan' and 'state' to numerical.
- Check for, and remove outliers

```
In [429]:  # drop unimportant columns  
data = data.drop(columns=['phone number', 'area code'], axis=1)  
  
# Confirm columns are dropped  
data.columns
```

```
Out[429]: Index(['state', 'account length', 'international plan', 'voice mail plan',  
               'number vmail messages', 'total day minutes', 'total day calls',  
               'total day charge', 'total eve minutes', 'total eve calls',  
               'total eve charge', 'total night minutes', 'total night calls',  
               'total night charge', 'total intl minutes', 'total intl calls',  
               'total intl charge', 'customer service calls', 'churn'],  
              dtype='object')
```

In [430]: ▶

```
# OneHotCode the three categorical columns of interest

data = pd.get_dummies(data, columns=['international plan','voice mail plan',

# Convert the one-hot encoded columns and the target colum 'Churn' from bo
for col in data.columns:
    if data[col].dtype == 'bool':
        data[col] = data[col].astype(int)

data.head()
```

Out[430]:

	account length	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes	total night calls	...
0	128	25	265.1	110	45.07	197.4	99	16.78	244.7	91	...
1	107	26	161.6	123	27.47	195.5	103	16.62	254.4	103	...
2	137	0	243.4	114	41.38	121.2	110	10.30	162.6	104	...
3	84	0	299.4	71	50.90	61.9	88	5.26	196.9	89	...
4	75	0	166.7	113	28.34	148.3	122	12.61	186.9	121	...

5 rows × 68 columns

◀

▶

In [431]: ▶ *# Preview the DataFrame*

```
data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 68 columns):
```

#	Column	Non-Null Count	Dtype
0	account length	3333 non-null	int64
1	number vmail messages	3333 non-null	int64
2	total day minutes	3333 non-null	float64
3	total day calls	3333 non-null	int64
4	total day charge	3333 non-null	float64
5	total eve minutes	3333 non-null	float64
6	total eve calls	3333 non-null	int64
7	total eve charge	3333 non-null	float64
8	total night minutes	3333 non-null	float64
9	total night calls	3333 non-null	int64
10	total night charge	3333 non-null	float64
11	total intl minutes	3333 non-null	float64
12	total intl calls	3333 non-null	int64
13	total intl charge	3333 non-null	float64
14	customer service calls	3333 non-null	int64
15	churn	3333 non-null	int32
16	international plan_yes	3333 non-null	uint8
17	voice mail plan_yes	3333 non-null	uint8
18	state_AL	3333 non-null	uint8
19	state_AR	3333 non-null	uint8
20	state_AZ	3333 non-null	uint8
21	state_CA	3333 non-null	uint8
22	state_CO	3333 non-null	uint8
23	state_CT	3333 non-null	uint8
24	state_DC	3333 non-null	uint8
25	state_DE	3333 non-null	uint8
26	state_FL	3333 non-null	uint8
27	state_GA	3333 non-null	uint8
28	state_HI	3333 non-null	uint8
29	state_IA	3333 non-null	uint8
30	state_ID	3333 non-null	uint8
31	state_IL	3333 non-null	uint8
32	state_IN	3333 non-null	uint8
33	state_KS	3333 non-null	uint8
34	state_KY	3333 non-null	uint8
35	state_LA	3333 non-null	uint8
36	state_MA	3333 non-null	uint8
37	state_MD	3333 non-null	uint8
38	state_ME	3333 non-null	uint8
39	state_MI	3333 non-null	uint8
40	state_MN	3333 non-null	uint8
41	state_MO	3333 non-null	uint8
42	state_MS	3333 non-null	uint8
43	state_MT	3333 non-null	uint8
44	state_NC	3333 non-null	uint8
45	state_ND	3333 non-null	uint8
46	state_NE	3333 non-null	uint8
47	state_NH	3333 non-null	uint8
48	state_NJ	3333 non-null	uint8
49	state_NM	3333 non-null	uint8
50	state_NV	3333 non-null	uint8
51	state_NY	3333 non-null	uint8

52	state_OH	3333	non-null	uint8
53	state_OK	3333	non-null	uint8
54	state_OR	3333	non-null	uint8
55	state_PA	3333	non-null	uint8
56	state_RI	3333	non-null	uint8
57	state_SC	3333	non-null	uint8
58	state_SD	3333	non-null	uint8
59	state_TN	3333	non-null	uint8
60	state_TX	3333	non-null	uint8
61	state_UT	3333	non-null	uint8
62	state_VA	3333	non-null	uint8
63	state_VT	3333	non-null	uint8
64	state_WA	3333	non-null	uint8
65	state_WI	3333	non-null	uint8
66	state_WV	3333	non-null	uint8
67	state_WY	3333	non-null	uint8

dtypes: float64(8), int32(1), int64(7), uint8(52)

memory usage: 573.0 KB

```
In [432]: def remove_outliers(data, columns):
            for col in columns:
                # Calculate Q1 (25th percentile) and Q3 (75th percentile)
                Q1 = data[col].quantile(0.25)
                Q3 = data[col].quantile(0.75)
                IQR = Q3 - Q1 # Interquartile Range

                # Define lower and upper bounds for detecting outliers
                lower_bound = Q1 - 1.5 * IQR
                upper_bound = Q3 + 1.5 * IQR

                # Filter out outliers
                data = data[(data[col] >= lower_bound) & (data[col] <= upper_bound)]

            return data

            # List of columns to check for outliers (excluding 'Churn')
            feature_columns = [col for col in data.columns if col != 'Churn' and data[col].nunique() > 1]

            # Apply the function to remove outliers
            df = remove_outliers(data, feature_columns)
            df.head()
```

Out[432]:

	account length	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes	total night calls	...
0	128	25	265.1	110	45.07	197.4	99	16.78	244.7	91	...
1	107	26	161.6	123	27.47	195.5	103	16.62	254.4	103	...
2	137	0	243.4	114	41.38	121.2	110	10.30	162.6	104	...
4	75	0	166.7	113	28.34	148.3	122	12.61	186.9	121	...
5	118	0	223.4	98	37.98	220.6	101	18.75	203.9	118	...

5 rows × 68 columns

## 2.2 Exploratory Data Analysis

```
In [433]: # create a copy of the clean dataframe
            df=df.copy(deep=True)
```

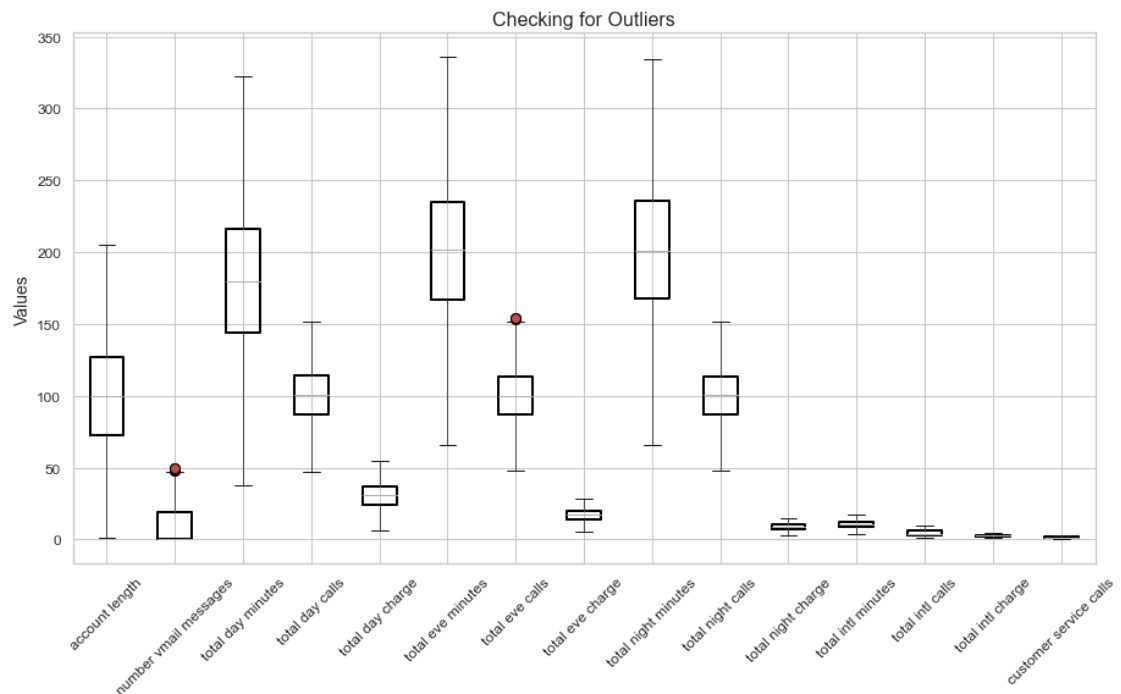
In [434]:

```

# Define a function to plot boxplots for cleaned columns
def plot_boxplots(df, columns):
    plt.figure(figsize=(15,8)) # Increase the figure size for better clarity
    df[columns].boxplot(boxprops=dict(linewidth=2), flierprops=dict(marker='o', color='red'))
    plt.title('Checking for Outliers', fontsize=16) # Add a more descriptive title
    plt.ylabel('Values', fontsize=14)
    plt.xticks(rotation=45, fontsize=12) # Rotate x-axis labels for better readability
    plt.yticks(fontsize=12) # Adjust y-axis label size for consistency
    plt.grid(True) # Add a grid for better visual reference
    plt.savefig('outliers')
    plt.show()

plot_boxplots(df, feature_columns)

```




## 2.2.1 Correlation Matrix

```
In [435]: # Calculate the correlation matrix
corr_matrix_columns = df[['account length',
    'number vmail messages', 'total day minutes', 'total day calls',
    'total day charge', 'total eve minutes', 'total eve calls',
    'total eve charge', 'total night minutes', 'total night calls',
    'total night charge', 'total intl minutes', 'total intl calls',
    'total intl charge', 'customer service calls', ]]
corr_matrix = corr_matrix_columns.corr().round(2)

plt.figure(figsize=(15, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix Heatmap')

plt.show()
```



- Proposed remedy is to drop one of the correlated predictors from the model as seen above in the heatmap there is clear correlation with an expectation of an impact on model performance and interpretability due to multicollinearity. There is a possibility of overfitting the model.

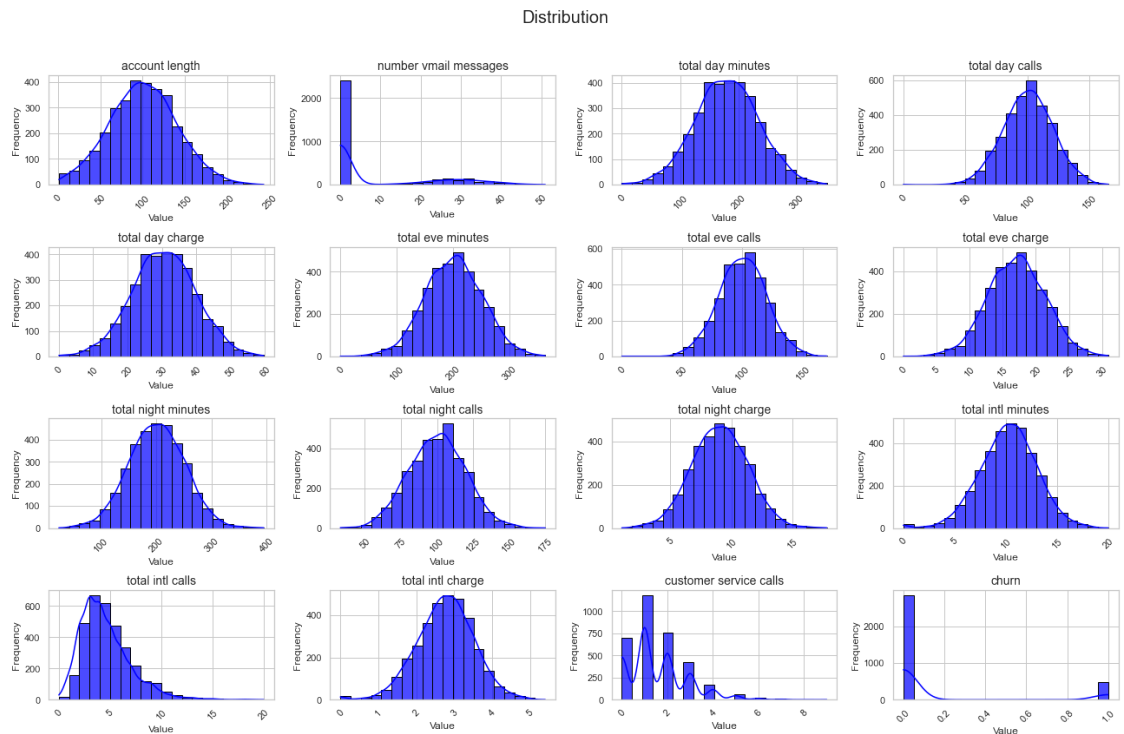
## 2.2.2 Histograms

```
In [436]: # List of continuous columns to plot
continuous_cols = df[['account length',
                      'number vmail messages', 'total day minutes', 'total day calls',
                      'total day charge', 'total eve minutes', 'total eve calls',
                      'total eve charge', 'total night minutes', 'total night calls',
                      'total night charge', 'total intl minutes', 'total intl calls',
                      'total intl charge', 'customer service calls', 'churn']]

# Create subplots
fig, axes = plt.subplots(nrows=4, ncols=4, figsize=(18, 12)) # Adjusted ;
fig.suptitle('Distribution', fontsize=20)

for i, col in enumerate(continuous_cols):
    ax = axes.flatten()[i]
    sns.histplot(data[col], bins=20, kde=True, color='blue', alpha=0.7, ec='black')
    ax.set_title(col, fontsize=14)
    ax.set_xlabel('Value', fontsize=12)
    ax.set_ylabel('Frequency', fontsize=12)
    ax.tick_params(axis='x', rotation=45) # Rotate x-axis labels for better readability

plt.tight_layout(rect=[0, 0, 1, 0.96]) # Adjust layout to accommodate the title
plt.savefig('Distribution')
plt.show()
```



- The charts show that the number of calls and charges for local calls are approximately normally distributed.
- The Customer Service Calls show distinct peaks at 1,2, indicating that these values are more frequent.

- Number of Voice Mail Messages is highly skewed to the right, with most values concentrated around 0 and a few concentrated at 20 - 40. Attributed to the service being an opt in service.

### 2.2.3 Class distribution of the target variable

```
In [437]: ▶ # Count the values of the churn column
churn_count = df['churn'].value_counts()

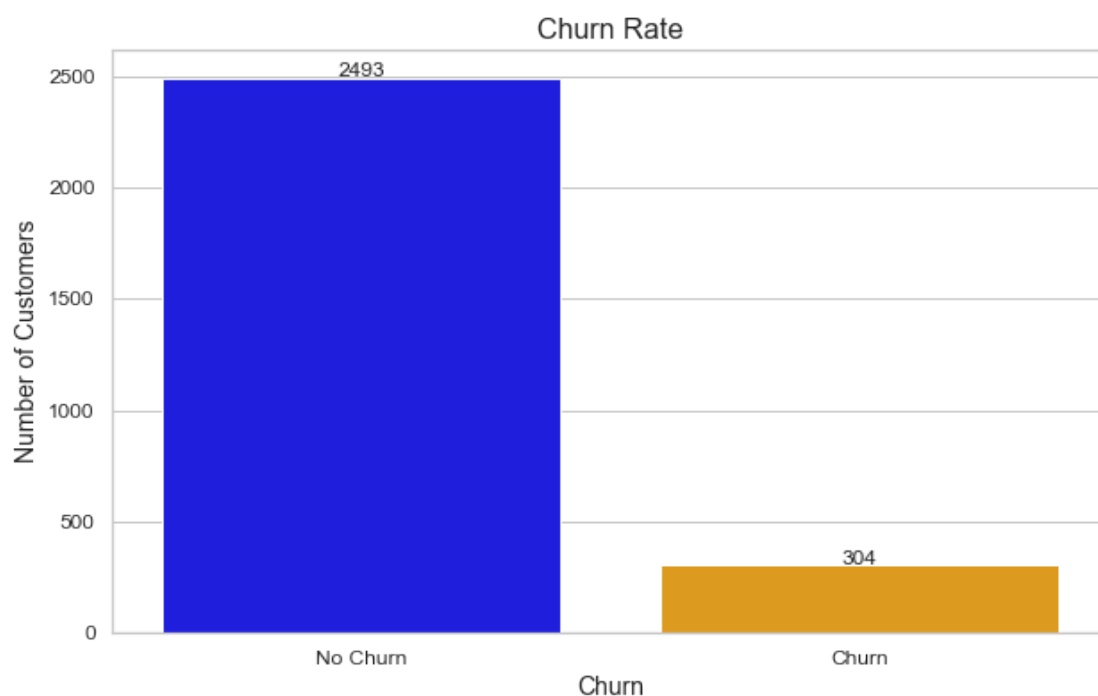
# Plot the bar chart
plt.figure(figsize=(10, 6))
sns.set(style="whitegrid")
ax = sns.barplot(x=churn_count.index, y=churn_count.values, palette=["blue", "orange"])

# Add titles and labels
plt.title('Churn Rate', fontsize=16)
plt.xlabel('Churn', fontsize=14)
plt.ylabel('Number of Customers', fontsize=14)
plt.xticks(ticks=[0, 1], labels=['No Churn', 'Churn'], fontsize=12)
plt.yticks(fontsize=12)

# Display the count values on top of the bars
for i, count in enumerate(churn_count.values):
    ax.text(i, count + 5, str(count), ha='center', fontsize=12)

plt.savefig('Churn_Rate')

plt.show()
```



From the above plot we can clearly come to a conclusion that there is a significant class imbalance. Therefore a significant impact on reliability of the model.

### 3.0 Modeling

I'll use a model iteration strategy in this part that takes advantage of feature importance and hyperparameter adjustment to address class imbalance.

### 3.1 Data Preparation

In [438]:

```
# In order to standardise the range of features to ensure they all contribute
from sklearn.preprocessing import MinMaxScaler # to scale the numeric features
transformer = MinMaxScaler()

def scaling(columns):
    return transformer.fit_transform(df[columns].values.reshape(-1,1))

for i in df.select_dtypes(include=[np.number]).columns:
    df[i] = scaling(i)
df.head()
```

Out[438]:

	account length	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total nigh minutes
0	0.622549	0.50	0.798455	0.600000	0.798430	0.486667	0.481132	0.486710	0.6654
1	0.519608	0.52	0.435042	0.723810	0.434944	0.479630	0.518868	0.479739	0.7014
2	0.666667	0.00	0.722261	0.638095	0.722222	0.204444	0.584906	0.204357	0.3602
4	0.362745	0.00	0.452949	0.628571	0.452912	0.304815	0.698113	0.305011	0.4505
5	0.573529	0.00	0.652037	0.485714	0.652003	0.572593	0.500000	0.572549	0.5137


5 rows × 68 columns

Remove Correlated Columns



```
In [439]: # Calculate the correlation matrix to remove multicollinearity
corr_matrix = df[['account length', 'number vmail messages', 'total day minutes',
                  'total eve minutes', 'total eve calls',
                  'total night minutes', 'total night calls',
                  'total intl minutes', 'total intl calls',
                  'customer service calls']]
corr_matrix = corr_matrix.corr().round(2)
corr_df = corr_matrix

plt.figure(figsize=(15, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix Heatmap after dropping columns')
plt.savefig('Correlation_Matrix_Heatmap')
plt.show()
```



### 3.2 Standardize the feature columns

```
In [440]: # predictors and Target represented by X and y

X = df.drop(columns=['churn','number vmail messages', 'total day charge',
                    'total intl charge'],axis=1)
y = df['churn']

# Split the data into a training and a test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,

# Intitalize the StandardScaler
scaler = StandardScaler()

# Fit the scaler on the training data and transform both the training and
X_train_standardized = scaler.fit_transform(X_train)
X_test_standardized = scaler.transform(X_test)

# Retain feature names and convert back to DataFrame
X_train = pd.DataFrame(X_train_standardized, columns=X_train.columns)
X_test = pd.DataFrame(X_test_standardized,columns=X_test.columns)
```

```
In [441]: # Check the shape of the standardized X_datasets
print(f"The y_train data set consists of {y_train.shape[0]} rows")
print(f"The X_train data set consists of {X_train.shape[0]} rows")
print(f"The X_train data set consists of {X_train.shape[1]} columns\n")

print(f"The y_test data set consists of {y_test.shape[0]} rows")
print(f"The X_test data set consists of {X_test.shape[0]} rows")
print(f"The X_train data set consists of {X_test.shape[1]} columns")
```

```
The y_train data set consists of 2097 rows
The X_train data set consists of 2097 rows
The X_train data set consists of 62 columns
```

```
The y_test data set consists of 700 rows
The X_test data set consists of 700 rows
The X_train data set consists of 62 columns
```

Both the training and test features have been standardized in order to make the model training and evaluation more reliable and effective. The two data sets have a 75.25 split.

### 3.3 Baseline Logistic Regression Model

```

In [442]: # Import the necessary Libraries
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, roc_curve, auc

# Instantiate LogisticRegression
logreg = LogisticRegression(fit_intercept=False, solver='liblinear', C=1e5)

# Fit to training data
base_log_model = logreg.fit(X_train, y_train)

# Predict on train and test sets
y_hat_train = logreg.predict(X_train)
y_hat_test = logreg.predict(X_test)

# Get Accuracy Score
print('Training Accuracy:', round(accuracy_score(y_train, y_hat_train), 2))
print('Test Accuracy:', round(accuracy_score(y_test, y_hat_test), 2))

# Create the ROC Curve for both the train and test sets

# Calculate the probability scores of each point for the train and test sets
y_train_score = base_log_model.decision_function(X_train)
y_test_score = base_log_model.decision_function(X_test)

# Calculate the fpr, tpr, and thresholds for the train and test sets
train_fpr, train_tpr, _ = roc_curve(y_train, y_train_score)
test_fpr, test_tpr, _ = roc_curve(y_test, y_test_score)

# Print the AUC for the train and test sets
print('Train AUC: {:.2f}'.format(auc(train_fpr, train_tpr)))
print('Test AUC: {:.2f}'.format(auc(test_fpr, test_tpr)))

# Plot the ROC curves for the train and test sets
plt.figure(figsize=(16,14))
lw = 2

plt.plot(train_fpr, train_tpr, color='green', lw=lw, label='Train ROC curve')
plt.plot(test_fpr, test_tpr, color='yellow', lw=lw, label='Test ROC curve')

plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.yticks([i/20.0 for i in range(21)])
plt.xticks([i/20.0 for i in range(21)])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (Baseline Model)')
plt.legend(loc='lower right')
plt.show()


```

Training Accuracy: 0.64

Test Accuracy: 0.57

Train AUC: 0.89

Test AUC: 0.82



```

In [443]: ▶ # Generate predictions using baseline_model and X_train
y_pred = logreg.predict(X_test)

# Ensure the length of y_test and y_pred are consistent
if len(y_test) != len(y_pred):
    raise ValueError(f"Inconsistent number of samples: y_test has {len(y_t

# Print classification metrics
print("***** LOGISTIC REGRESSION CLASSIFIER MODEL RESULTS *****")
print(classification_report(y_test, y_pred, target_names=['No Churn', 'Ch

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

# Plot Confusion Matrix
plt.figure(figsize=(8, 6)) # Adjust the figure size for better clarity
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Churn
plt.xlabel('Predicted', fontsize=12)
plt.ylabel('Actual', fontsize=12)
plt.title('Confusion Matrix', fontsize=15)
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.tight_layout()
plt.show()

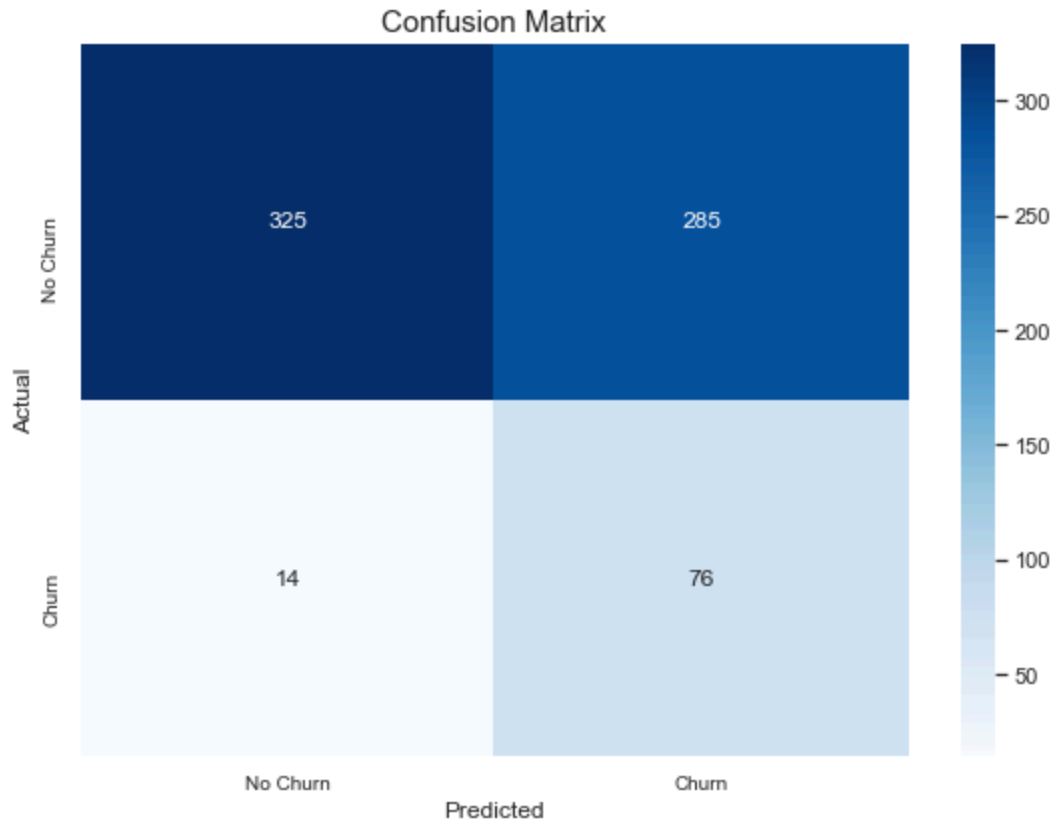
```

```

***** LOGISTIC REGRESSION CLASSIFIER MODEL RESULTS *****
*****

```

	precision	recall	f1-score	support
No Churn	0.96	0.53	0.68	610
Churn	0.21	0.84	0.34	90
accuracy			0.57	700
macro avg	0.58	0.69	0.51	700
weighted avg	0.86	0.57	0.64	700



The confusion matrix above indicates the performance of the model as below:

- True positive; There are 76 instances of the churn being correctly predicted as "Churn"
- True Negative; There are 325 instances of the churn correctly predicted as "No Churn"
- False Positive; There are 285 instances of the churn being predicted as "Churn" where its actually "No Churn"
- False Negative; There are 14 instances of the churn being predicted as "No Churn" where its actually "Churn"

Overall, the model does a good job of detecting "No Churn" cases, but it has a lot of struggle detecting "Churn," as shown by the large amount of false negatives. This implies that there may be an imbalance in the dataset or that the model's sensitivity to the "Churn" class should be improved.

```
In [444]: ▶ # Calculate and print key metrics
print("***** LOGISTIC REGRESSION CLASSIFIER MODEL RESULTS *****")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.5f}")
print(f"Precision: {precision_score(y_test, y_pred):.5f}")
print(f"Recall: {recall_score(y_test, y_pred):.5f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.5f}")

***** LOGISTIC REGRESSION CLASSIFIER MODEL RESULTS *****
*****
Accuracy: 0.57286
Precision: 0.21053
Recall: 0.84444
F1 Score: 0.33703
```

The Logistics Regressionmodel performance metrics are as follows:

- Accuracy ( 0.57286): The model accurately predicted 57% of all instances.
- Precision (0.21053): Of the cases predicted as "Churn," 21% were correct.
- Recall (0.84444): The model successfully identified only 84% of the actual "Churn" cases.
- F1 Score (0.33703): The low F1 score reflects poor overall performance in detecting "Churn," balancing both precision and recall.

These metrics are particularly useful for imbalanced datasets, as Accuracy alone may not reflect the model's ability to correctly identify the minority class ("Churn"). In this case, the metrics will highlight that while the model performs well in predicting "No Churn," it has lower Recall and F1 Score for "Churn," indicating room for improvement in recognizing this minority class.

### 3.4 Random Forest model

```
In [445]: # Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_sampled, y_sampled = smote.fit_resample(X_train, y_train)

# Train the Random Forest model
model = RandomForestClassifier(random_state=42)
model.fit(X_sampled, y_sampled)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate performance
print("***** RANDOM FOREST CLASSIFIER RESULTS *****")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.5f}")
print(f"Precision: {precision_score(y_test, y_pred):.5f}")
print(f"Recall: {recall_score(y_test, y_pred):.5f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.5f}")

***** RANDOM FOREST CLASSIFIER RESULTS *****
Accuracy: 0.91143
Precision: 0.72581
Recall: 0.50000
F1 Score: 0.59211
```

The Random Forest model which has SMOTE applied to it clearly outperforms the Logistic Regression baseline in all metrics. While Logistic Regression shows acceptable accuracy, its poor recall and F1 score highlight its inability to effectively detect "Churn." In contrast, Random Forest demonstrates strong performance across all metrics, making it a much better choice for this problem, especially if identifying "Churn" is critical.



```
In [446]: # Handle class imbalance with SMOTE
smote = SMOTE(random_state=42)
X_sampled, y_sampled = smote.fit_resample(X_train, y_train)

# Define the Random Forest model
rf = RandomForestClassifier(random_state=42)

# Define the hyperparameters grid
param_grid = {
    'n_estimators': [50, 100, 200],           # Number of trees in the forest
    'max_depth': [None, 10, 20, 30],          # Maximum depth of the tree
    'min_samples_split': [2, 5, 10],           # Minimum number of samples
    'min_samples_leaf': [1, 2, 4],             # Minimum number of samples
    'bootstrap': [True, False]                # Whether bootstrap samples
}

# Apply GridSearchCV
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_sampled, y_sampled)

# Best hyperparameters and model
best_rf = grid_search.best_estimator_
print("Best Hyperparameters:", grid_search.best_params_)

# Make predictions with the best model
y_pred = best_rf.predict(X_test)

# Evaluate the tuned model
print("***** TUNED RANDOM FOREST CLASSIFIER RESULTS *****")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.5f}")
print(f"Precision: {precision_score(y_test, y_pred):.5f}")
print(f"Recall: {recall_score(y_test, y_pred):.5f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.5f}")
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 17 tasks      | elapsed:    8.5s
[Parallel(n_jobs=-1)]: Done 138 tasks     | elapsed:   25.5s
[Parallel(n_jobs=-1)]: Done 341 tasks     | elapsed:   52.6s
[Parallel(n_jobs=-1)]: Done 624 tasks     | elapsed:   1.6min
[Parallel(n_jobs=-1)]: Done 989 tasks     | elapsed:   2.4min
[Parallel(n_jobs=-1)]: Done 1080 out of 1080 | elapsed:   2.6min finished
```

Best Hyperparameters: {'bootstrap': False, 'max\_depth': 30, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}

\*\*\*\*\* TUNED RANDOM FOREST CLASSIFIER RESULTS \*\*\*\*\*

Accuracy: 0.90571

Precision: 0.73077

Recall: 0.42222

F1 Score: 0.53521

In [447]:

```

# Create a DataFrame to hold feature importances
importance_df = pd.DataFrame({'Feature': X_train.columns, 'Importance': fe

# Sort the DataFrame by importance and select the top 10 features
top_features_df = importance_df.sort_values(by='Importance', ascending=False)

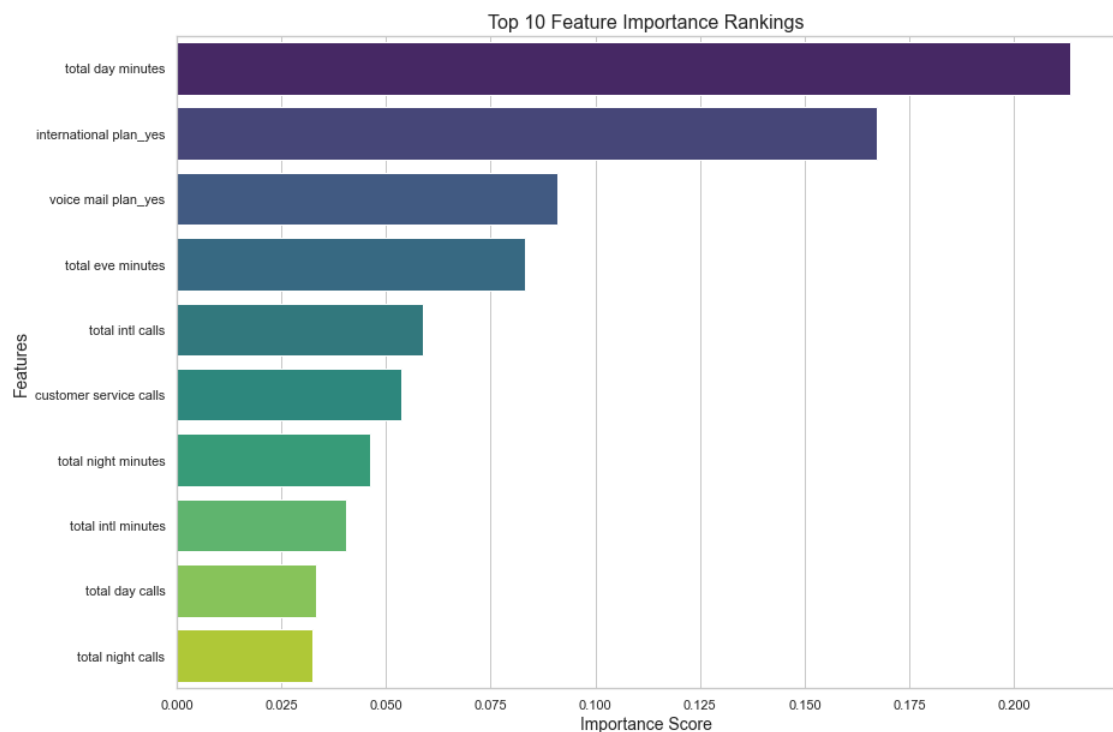
# Set the plot style
sns.set(style="whitegrid")

# Create a bar plot for the top 10 features
plt.figure(figsize=(14, 10))
sns.barplot(x='Importance', y='Feature', data=top_features_df, palette="v

# Add title and labels
plt.title('Top 10 Feature Importance Rankings', fontsize=16)
plt.xlabel('Importance Score', fontsize=14)
plt.ylabel('Features', fontsize=14)

# Show the plot & Save the plot
plt.savefig('Top_10_Feature_importance_Rankings')
plt.show()

```



### 3.5 Decision Tree

```
In [448]: ▶ # Handle class imbalance with SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

# Define the Decision Tree model
dt = DecisionTreeClassifier(random_state=42)

# Define the hyperparameter space for Randomized Search
param_dist = {
    'criterion': ['gini', 'entropy'],           # Splitting criterion
    'max_depth': [None, 10, 20, 30, 50],        # Maximum depth of the tree
    'min_samples_split': [2, 5, 10, 20],         # Minimum samples to split
    'min_samples_leaf': [1, 2, 4, 10],          # Minimum samples at a leaf
    'max_features': [None, 'sqrt', 'log2'],     # Number of features to consider
    'splitter': ['best', 'random']              # Strategy for choosing splits
}

# Apply RandomizedSearchCV
random_search = RandomizedSearchCV(
    estimator=dt,
    param_distributions=param_dist,
    n_iter=100,                               # Number of parameter settings sampled
    scoring='f1',                             # Use F1 score to evaluate performance
    cv=5,                                     # 5-fold cross-validation
    random_state=42,                           # Ensures reproducibility
    verbose=2,
    n_jobs=-1                                 # Use all available processors
)
random_search.fit(X_resampled, y_resampled)

# Best hyperparameters and model
best_dt = random_search.best_estimator_
print("Best Hyperparameters:", random_search.best_params_)

# Make predictions with the best model
y_pred = best_dt.predict(X_test)

# Step 7: Evaluate the tuned model
print("***** TUNED DECISION TREE RESULTS *****")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.5f}")
print(f"Precision: {precision_score(y_test, y_pred):.5f}")
print(f"Recall: {recall_score(y_test, y_pred):.5f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.5f}")
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 17 tasks | elapsed: 0.1s

[Parallel(n\_jobs=-1)]: Done 426 tasks | elapsed: 1.8s

[Parallel(n\_jobs=-1)]: Done 477 out of 500 | elapsed: 2.0s remaining: 0.0s

```
Best Hyperparameters: {'splitter': 'best', 'min_samples_split': 2, 'min_samples_leaf': 2, 'max_features': None, 'max_depth': 20, 'criterion': 'entropy'}
```

```
***** TUNED DECISION TREE RESULTS *****
```

```
Accuracy: 0.89571
```

```
Precision: 0.58416
```

```
Recall: 0.65556
```

```
F1 Score: 0.61780
```

```
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 2.1s finished
```

The Decision Tree has lower accuracy and precision compared to Random Forest, with a higher recall. It's prone to overfitting due to its single tree structure, which could explain the imbalance between precision and recall. The Random Forest performs well across all metrics, with high accuracy and reasonable precision and recall. Its ensemble nature (using multiple trees) helps in reducing overfitting and improving stability.



```

In [449]: # Define classifiers
classifiers = [
    LogisticRegression(max_iter=1000), # Increased max_iter for Logistic
    RandomForestClassifier(),
    DecisionTreeClassifier()
]

# Define result tables for training and test data
train_result_table = pd.DataFrame(columns=['classifiers', 'auc', 'accuracy'])
test_result_table = pd.DataFrame(columns=['classifiers', 'auc', 'accuracy'])

# Train the models and record the results
for cls in classifiers:
    model = cls.fit(X_train, y_train)

    # Training data predictions
    y_train_proba = model.predict_proba(X_train)[: , 1]
    y_train_pred = model.predict(X_train)
    train_auc = roc_auc_score(y_train, y_train_proba)
    train_accuracy = accuracy_score(y_train, y_train_pred)

    # Test data predictions
    y_test_proba = model.predict_proba(X_test)[: , 1]
    y_test_pred = model.predict(X_test)
    test_auc = roc_auc_score(y_test, y_test_proba)
    test_accuracy = accuracy_score(y_test, y_test_pred)

    # Append results for training and test data
    train_result_table = pd.concat([train_result_table,
                                    pd.DataFrame({'classifiers': [cls.__name__],
                                                    'auc': [train_auc],
                                                    'accuracy': [train_accuracy]})],
                                   ignore_index=True)
    test_result_table = pd.concat([test_result_table,
                                   pd.DataFrame({'classifiers': [cls.__name__],
                                                  'auc': [test_auc],
                                                  'accuracy': [test_accuracy]})],
                                  ignore_index=True)

# Identify the best model for training and test data
best_train_model = train_result_table.loc[train_result_table['auc'].idxmax()]
best_test_model = test_result_table.loc[test_result_table['auc'].idxmax()]

# Display comparison results
print("***** MODEL COMPARISON RESULTS *****")
print("Training Data:")
print(train_result_table)
print("\nBest Model on Training Data: {} (AUC: {:.3f}, Accuracy: {:.3f})".format(
    best_train_model['classifiers'], best_train_model['auc'], best_train_model['accuracy']))

print("\nTest Data:")
print(test_result_table)
print("\nBest Model on Test Data: {} (AUC: {:.3f}, Accuracy: {:.3f})".format(
    best_test_model['classifiers'], best_test_model['auc'], best_test_model['accuracy']))

# Plot ROC curves for training and test data
plt.figure(figsize=(18, 12))

```

```

# Training ROC curves
plt.subplot(1, 2, 1)
plt.title('ROC Curve Analysis (Training Data)', fontweight='bold', fontsize=15)
for cls in classifiers:
    model = cls.fit(X_train, y_train)
    y_train_proba = model.predict_proba(X_train)[: , 1]
    fpr, tpr, _ = roc_curve(y_train, y_train_proba)
    auc = roc_auc_score(y_train, y_train_proba)
    plt.plot(fpr, tpr, label="{}, AUC={:.3f}".format(cls.__class__.__name__, auc))
plt.plot([0, 1], [0, 1], color='orange', linestyle='--')
plt.xlabel("False Positive Rate", fontsize=12)
plt.ylabel("True Positive Rate", fontsize=12)
plt.legend(prop={'size': 10}, loc='lower right')

# Test ROC curves
plt.subplot(1, 2, 2)
plt.title('ROC Curve Analysis (Test Data)', fontweight='bold', fontsize=15)
for cls in classifiers:
    model = cls.fit(X_train, y_train)
    y_test_proba = model.predict_proba(X_test)[: , 1]
    fpr, tpr, _ = roc_curve(y_test, y_test_proba)
    auc = roc_auc_score(y_test, y_test_proba)
    plt.plot(fpr, tpr, label="{}, AUC={:.3f}".format(cls.__class__.__name__, auc))
plt.plot([0, 1], [0, 1], color='orange', linestyle='--')
plt.xlabel("False Positive Rate", fontsize=12)
plt.ylabel("True Positive Rate", fontsize=12)
plt.legend(prop={'size': 10}, loc='lower right')

plt.tight_layout()
plt.savefig('Roc_Curve_Training_and_Test_Data')
plt.show()

```

\*\*\*\*\* MODEL COMPARISON RESULTS \*\*\*\*\*

Training Data:

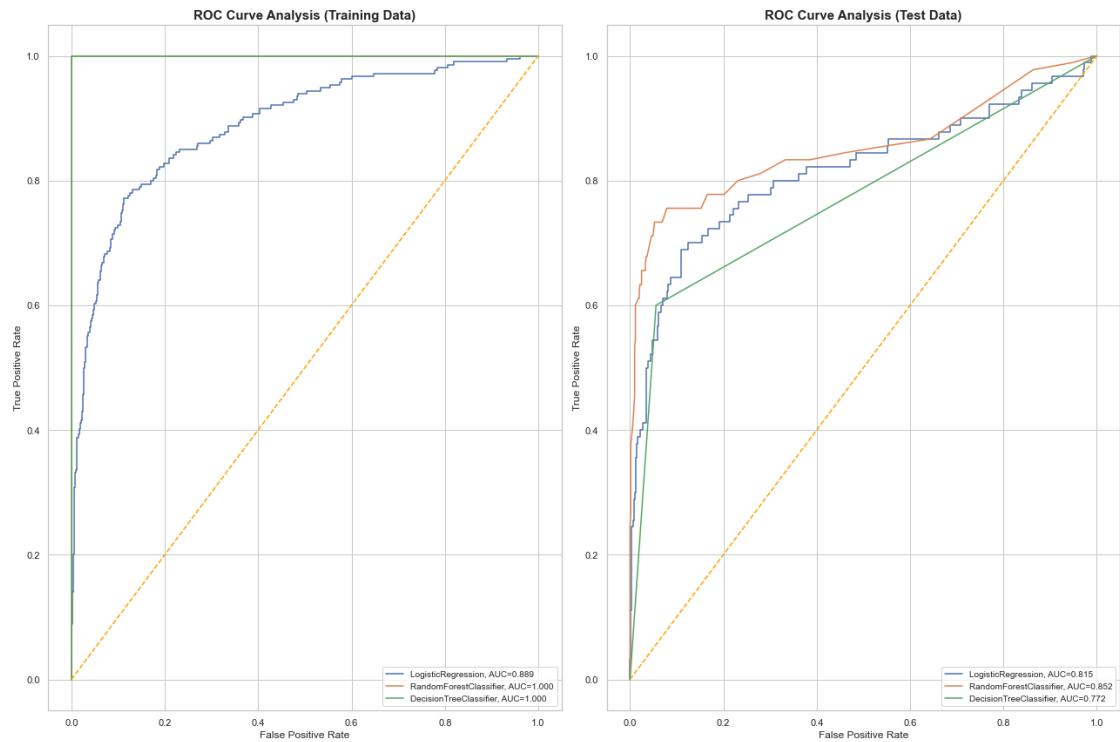
	classifiers	auc	accuracy
0	LogisticRegression	0.888664	0.925608
1	RandomForestClassifier	1.000000	1.000000
2	DecisionTreeClassifier	1.000000	1.000000

Best Model on Training Data: RandomForestClassifier (AUC: 1.000, Accuracy: 1.000)

Test Data:

	classifiers	auc	accuracy
0	LogisticRegression	0.814517	0.905714
1	RandomForestClassifier	0.852122	0.912857
2	DecisionTreeClassifier	0.782605	0.910000

Best Model on Test Data: RandomForestClassifier (AUC: 0.852, Accuracy: 0.913)



Of the models based on their AUC and accuracy scores for both training and test data we can conclude as follows;

- Accuracy is a measure of how often the model gets the prediction right, and in this case measures how often the model correctly predicts whether a customer will churn or not. An test accuracy score of 0.91 from the best Random Forest Ensemble Model means that our model was able to predict correctly 91% of the time.
- Random Forest and Decision Tree models have perfect accuracy and AUC scores of 1.000, suggesting they fit the training data perfectly. However, this could indicate overfitting, as these models may have memorized the training data rather than generalizing well.
- When evaluated on the test data, the Random Forest classifier stands out with the highest AUC (0.863270) and accuracy (92%). It outperforms the other two models, indicating better generalization and performance on unseen data. The Decision Tree model has a relatively high accuracy but a lower AUC, suggesting it may not handle the complexity of the data as well as Random Forest. Logistic Regression also has a lower AUC and accuracy compared to Random Forest.

**Conclusion:** The Random Forest classifier is the best model to use, as it achieves the highest accuracy and AUC on both training and test data, with the best generalization capability to unseen data.

## Recomendation and proposed customer retention strategies

- Focus on the most impactful features identified above to reduce the complexity of the model and make it more effective and effecient.



- Provide targeted offers and discounts to customers based on their patterns and preferences enhancing retention.
- Analyze customer service interaction to identify common ailing issues to the customers and address each promptly.
- address any issue that may hinder experience enhancement.